

CS5841/EE5841 Machine Learning

Lecture 3: Regression

Evan Lucas



Michigan Tech

Overview

- Course updates
-




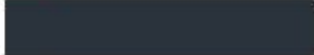


Class updates

- Julia module created
 - Extra credit quiz
 - Contains download links for Julia and Pluto
- Discord and discussion threads made



Course survey results

What is your experience with Python?

Expert - I have contributed to multiple large Python projects.	8 respondents	10 %	
Intermediate - somewhere between a beginner and an expert.	67 respondents	82 %	
Beginner - I know basic syntax.	7 respondents	9 %	
Why are you asking about snakes in a Machine Learning course?		0 %	



Course survey results

Expert - I have contributed to multiple large Julia projects.		0 %	✓
Intermediate - somewhere between a beginner and an expert.	2 respondents	2 %	
Beginner - I know basic syntax.	17 respondents	21 %	
I have no idea who Julia is in the context of Machine Learning	63 respondents	77 %	

This is a great idea. Let's learn Julia!	68 respondents	83 %	✓
This is a terrible idea. I only want to use Python.	12 respondents	15 %	
I don't care.	2 respondents	2 %	



Course survey results

Yes to Discord	26 respondents	32 %	<div><div></div></div> ✓
Yes to Discussions forum on Canvas	11 respondents	13 %	<div><div></div></div>
Yes to both	45 respondents	55 %	<div><div></div></div>
Yes to something else (will email you at eglucas@mtu.edu)		0 %	<div><div></div></div>
No. I only talk in person with people or on private channels that I control		0 %	<div><div></div></div>



Bonus topic question

- Lots of interest in:
 - Ethics
 - Large language models
 - NLP in general
 - Object detection
- Lecture plan to come soon.
 - Do you prefer a high level survey of a few topics or one deeper topic?
 - Do you want an extra credit assignment?
- A couple answers were generated with ChatGPT or similar...
 - Please don't do that unless you attribute it



Related reading

- Strongly suggested
 - Bishop Chapter 3
- Additional
 - Chapter 11 in Murphy
 - Chapter 3 in ESLII



Where does linear regression fit in the world of ML?

- Supervised learning
 - We know the answer and are training the model to predict it
- Regression
 - We are predicting a numerical value



Supervised learning formal definition

- Labeled datasets are used to train an algorithm to predict an outcome
- Given



Regression vs. Classification

- Regression
 - We are trying to predict a continuous value
- Classification
 - We are trying to predict a discrete value



Why linear regression?

- Simple models are useful
 - Always good to benchmark with a simple model!
 - Simple models are good placeholders in system backend prototyping
- Interpretable!
- A component of neural networks!
 - Called a linear, dense, or fully connected layer
- Introduction of concepts without a complicated model
 - Loss metrics
 - Probabilistic model training (gradient descent family)
 - Regularization



Linear regression

- $y = m \cdot x + b$
 - Given two points, can find an exact solution for a line
 - How do we handle multiple data points?
- $\{y\} = [w][x]$
 - Matrix definition
 - x_0 is 1




Some notes on notation

- Instead of algebraic standard form ($x^2+x+\dots$), we typically list things in ascending order
 - ie: w_0, w_1, \dots, w_n
 - This sets up the bias term as being the first column (or row in some references)



Solving linear regression

- Minimize difference between estimate and true value by adjusting weights
 - $\operatorname{argmin}_w (y - y_{\text{est}})^2$
 - $\operatorname{argmin}_w (y - wx)^2$
- Why squared error?
 - Always positive
 - Easy to differentiate



Solving linear regression

$$\arg \min_w \sum_i (y_i - wx_i)^2$$

$$\frac{\delta}{\delta w} \sum_i (y_i - wx_i)^2 = 2 \sum_i -x_i(y_i - wx_i)$$


$$2 \sum_i x_i(y_i - wx_i) = 0$$

$$2 \sum_i x_i y_i - 2 \sum_i wx_i x_i = 0$$

$$w = \frac{\sum_i x_i y_i}{\sum_i x_i^2}$$




Adding a bias term



$$E = \sum_i (y_i - w_0 - w_1 x_i)^2$$



Adding a bias term



$$w_0 = \bar{y} - w_1 \bar{x}$$

$$w_1 = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sum_i (x_i - \bar{x})^2}$$



Multiple linear regression

$$y = w_0 + w_1x_1 + \dots + w_kx_k$$



Basis functions

- We can use the concept of basis functions to map non-linear spaces into our linear model


$$y(\mathbf{x}, \mathbf{w}) = \sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})$$

- ϕ_j is our basis function
 - $\phi_d(x) = x_d$ is the linear basis function
 - $\phi_0(x) = 1$, typically, to give us a bias term



Polynomial basis functions


- Global!
 - Small change in \mathbf{x} affects all basis functions



$$\phi_j(x) = x^j$$




Example of polynomial curve fitting




$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$




0th Order Polynomial




1st Order Polynomial




3rd Order Polynomial



9th Order Polynomial




Concept: Overfitting!



Gaussian basis functions

- Local!
 - Small changes in \mathbf{x} only affect nearby basis functions
 - Parameters control location and scale (width)




$$\phi_j(x) = \exp \left\{ -\frac{(x - \mu_j)^2}{2s^2} \right\}$$



Sigmoidal basis functions

- Local!
- Scale parameter affects slope



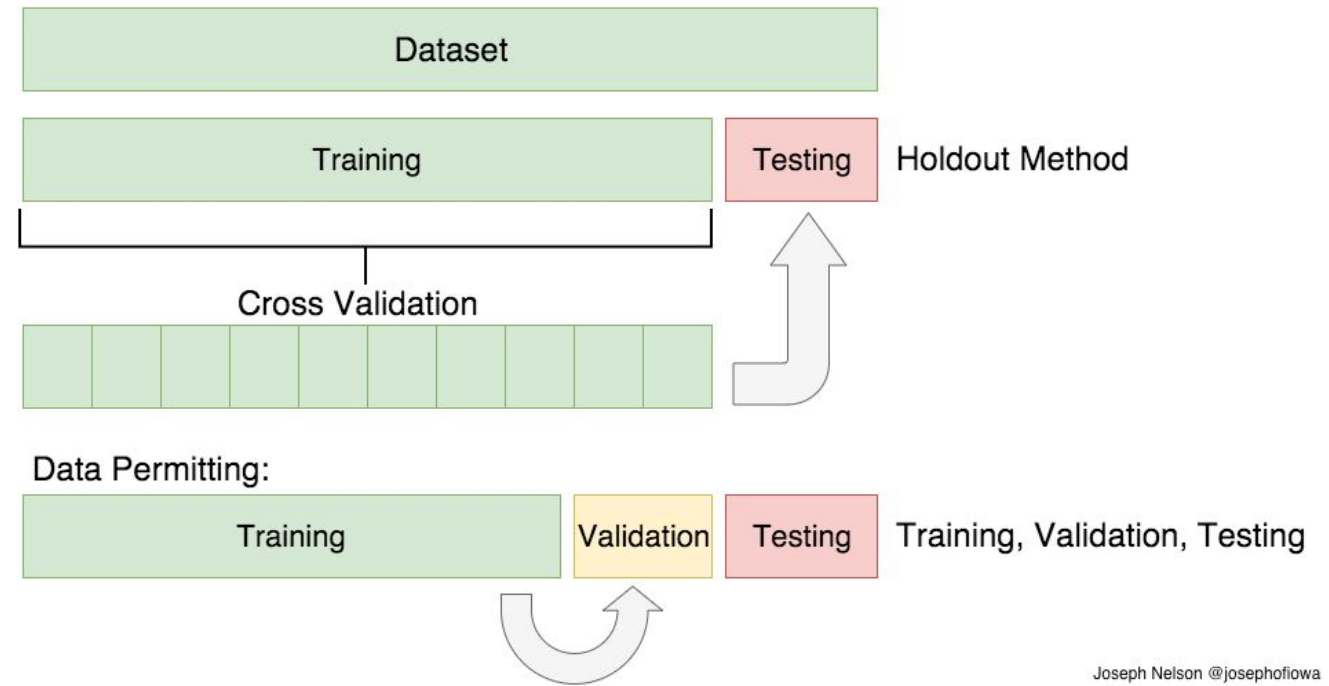
$$\phi_j(x) = \sigma\left(\frac{x - \mu_j}{s}\right)$$

$$\sigma(a) = \frac{1}{1 + \exp(-a)}.$$



Data splits

- We want a fair assessment of our model!
 - This requires testing on data we didn't use for training
- Different strategies for test/train split
 - Test/train
 - Test/train/validation
 - Cross validation with separate test split



Joseph Nelson @josephofiowa

Test/train/validation splits

- Many big datasets contain test/train/validation splits
 - Some are just test/train
 - Common test split ensures fair comparison
 - Long training time precludes cross validation
- Split uses
 - Test - used to measure final model performance (usually 10-20%)
 - Train - used to train model (usually 70-90%)
 - Validation - used to evaluate model for model tuning (usually 10-20%)



Cross validation

- *K-fold cross validation*
 - Partition data X_{train} into K separate sets of equal size
 - $X_{\text{train}} = (X_{\text{train},1}, X_{\text{train},2}, \dots, X_{\text{train},K})$
 - Common K are K=5 and K=10
- For each $k=1,2,\dots,K$
 - Fit the model $y(w,\lambda)$ to the training set excluding kth fold $X_{\text{train},k}$
 - Compute values for $X_{\text{train},k}$ and compute error
 - Repeat for each



Special cases of cross validation

- What if we do n -fold cross validation, where n is the dataset size?
 - *Leave one out cross validation (LOO CV)*
 - Not very data efficient
 - Good for estimating model performance with all available data
 - Bad for estimating model generalization with unseen data



Data leakage

- Extra data is available during training that is not available during testing/inference
- Ex: this paper (Learning with Signatures [1]) scored 100% on several common benchmark datasets
 - Why is this suspicious?
 - How did they do it?
 - They created different classifiers for each class and only used that classifier for that class!
- Simpler example: predicting yearly salary and including a monthly_salary variable

[1] <https://arxiv.org/abs/2204.07953v1>



How does data leakage happen?

- Improper featurization
 - Data pre-processing that shares information is fit on test and train splits instead of just test split
- Duplicate datapoints
 - Easy to do when oversampling or augmenting
- Group leakage
 - Ex: dataset includes 1000 patients with 10 x-rays from each, not splitting by patient could cause model to learn patient instead of pathology
- Time leakage
 - Time series data is especially challenging!
 - Generally (not always), you split with older data in training and new data in test



Following slides are directly from Bishop Ch 3

Derivation will be gone over in class, but extra slides are here for your reference. The Bishop ch. 3 should also be a good reference



- Assume observations from a deterministic function with added Gaussian noise:

$$t = y(\mathbf{x}, \mathbf{w}) + \epsilon \quad \text{where} \quad p(\epsilon|\beta) = \mathcal{N}(\epsilon|0, \beta^{-1})$$

- which is the same as saying,

$$p(t|\mathbf{x}, \mathbf{w}, \beta) = \mathcal{N}(t|y(\mathbf{x}, \mathbf{w}), \beta^{-1}).$$

- Given observed inputs, $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, and targets, $\mathbf{t} = [t_1, \dots, t_N]^T$, we obtain the likelihood function

$$p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n|\mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}).$$

- Taking the logarithm, we get

$$\begin{aligned}\ln p(\mathbf{t}|\mathbf{w}, \beta) &= \sum_{n=1}^N \ln \mathcal{N}(t_n | \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n), \beta^{-1}) \\ &= \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi) - \beta E_D(\mathbf{w})\end{aligned}$$

- where

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n)\}^2$$

- is the sum-of-squares error.

- Computing the gradient and setting it to zero yields

$$\nabla_{\mathbf{w}} \ln p(\mathbf{t}|\mathbf{w}, \beta) = \beta \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\} \phi(\mathbf{x}_n)^T = \mathbf{0}.$$

- Solving for \mathbf{w} , we get

$$\mathbf{w}_{\text{ML}} = \left(\Phi^T \Phi \right)^{-1} \Phi^T \mathbf{t}$$

The Moore-Penrose pseudo-inverse, Φ^\dagger .

- where

$$\Phi = \begin{pmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \cdots & \phi_{M-1}(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & \phi_1(\mathbf{x}_2) & \cdots & \phi_{M-1}(\mathbf{x}_2) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}_N) & \phi_1(\mathbf{x}_N) & \cdots & \phi_{M-1}(\mathbf{x}_N) \end{pmatrix}.$$

- Maximizing with respect to the bias, w_0 , alone, we see that

$$w_0 = \bar{t} - \sum_{j=1}^{M-1} w_j \overline{\phi_j}$$

We can also maximize with respect to the noise precision parameter, \square


$$\frac{1}{\beta_{\text{ML}}} = \frac{1}{N} \sum_{n=1}^N \{t_n - \mathbf{w}_{\text{ML}}^T \boldsymbol{\phi}(\mathbf{x}_n)\}^2$$

- Data items considered one at a time (a.k.a. online learning); use stochastic (sequential) gradient descent:

$$\begin{aligned}\mathbf{w}^{(\tau+1)} &= \mathbf{w}^{(\tau)} - \eta \nabla E_n \\ &= \mathbf{w}^{(\tau)} + \eta (t_n - \mathbf{w}^{(\tau)\top} \phi(\mathbf{x}_n)) \phi(\mathbf{x}_n).\end{aligned}$$

- This is known as the *least-mean-squares (LMS) algorithm*. Issue: how to choose η ?

Concept: Overfitting!



- Consider the error function:

$$E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$$

Data term + Regularization term

- With the sum-of-squares error function and a quadratic regularizer, we get

$$\frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$



- which is minimized by

$$\mathbf{w} = \left(\lambda \mathbf{I} + \Phi^T \Phi \right)^{-1} \Phi^T \mathbf{t}.$$


λ is called the regularization coefficient.

- With a more general regularizer, we have


$$\frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^q$$




Lasso



Quadratic



Lasso tends to generate sparser solutions than a quadratic regularizer.



- Recall the *expected squared loss*,

$$\mathbb{E}[L] = \int \{y(\mathbf{x}) - h(\mathbf{x})\}^2 p(\mathbf{x}) d\mathbf{x} + \underbrace{\iint \{h(\mathbf{x}) - t\}^2 p(\mathbf{x}, t) d\mathbf{x} dt}_{\text{noise term}}$$

- where

$$h(\mathbf{x}) = \mathbb{E}[t|\mathbf{x}] = \int t p(t|\mathbf{x}) dt.$$

- The second term of $\mathbf{E}[L]$ corresponds to the noise inherent in the random variable t .
- What about the first term?

- Suppose we were given multiple data sets, each of size N . Any particular data set, \mathcal{D} , will give a particular function $y(\mathbf{x}; \mathcal{D})$. We then have

$$\begin{aligned} & \{y(\mathbf{x}; \mathcal{D}) - h(\mathbf{x})\}^2 \\ &= \{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] + \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}^2 \\ &= \{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})]\}^2 + \{\mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}^2 \\ &\quad + 2\{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})]\}\{\mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}. \end{aligned}$$

- Taking the expectation over \mathcal{D} yields

$$\begin{aligned} & \mathbb{E}_{\mathcal{D}} [\{y(\mathbf{x}; \mathcal{D}) - h(\mathbf{x})\}^2] \\ &= \underbrace{\{\mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}^2}_{(\text{bias})^2} + \underbrace{\mathbb{E}_{\mathcal{D}} [\{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})]\}^2]}_{\text{variance}}. \end{aligned}$$

- Thus we can write

$$\text{expected loss} = (\text{bias})^2 + \text{variance} + \text{noise}$$


- where

$$(\text{bias})^2 = \int \{\mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})] - h(\mathbf{x})\}^2 p(\mathbf{x}) \, d\mathbf{x}$$

$$\text{variance} = \int \mathbb{E}_{\mathcal{D}} [\{y(\mathbf{x}; \mathcal{D}) - \mathbb{E}_{\mathcal{D}}[y(\mathbf{x}; \mathcal{D})]\}^2] p(\mathbf{x}) \, d\mathbf{x}$$

$$\text{noise} = \iint \{h(\mathbf{x}) - t\}^2 p(\mathbf{x}, t) \, d\mathbf{x} \, dt$$

From these plots, we note that an over-regularized model (large λ) will have a high bias, while an under-regularized model (small λ) will have a high variance.



Questions + Comments?

