

# CS5841/EE5841 Machine Learning

## Lecture 6: Clustering and dimensionality reduction

Evan Lucas



Michigan Tech

# Overview

- Course updates
- Dimensionality reduction
- Clustering



# Class updates

- HW1 due soon
- HW2 progress?
- HW3 released last Friday
- Clustering/dimensionality reduction quiz instead of homework?
- Decision tree/ensemble learning quiz coming up
- Probability quiz deadlines extended
- Review day for half of class on Wednesday



# Dimensionality Reduction



Bishop 12



Michigan Tech

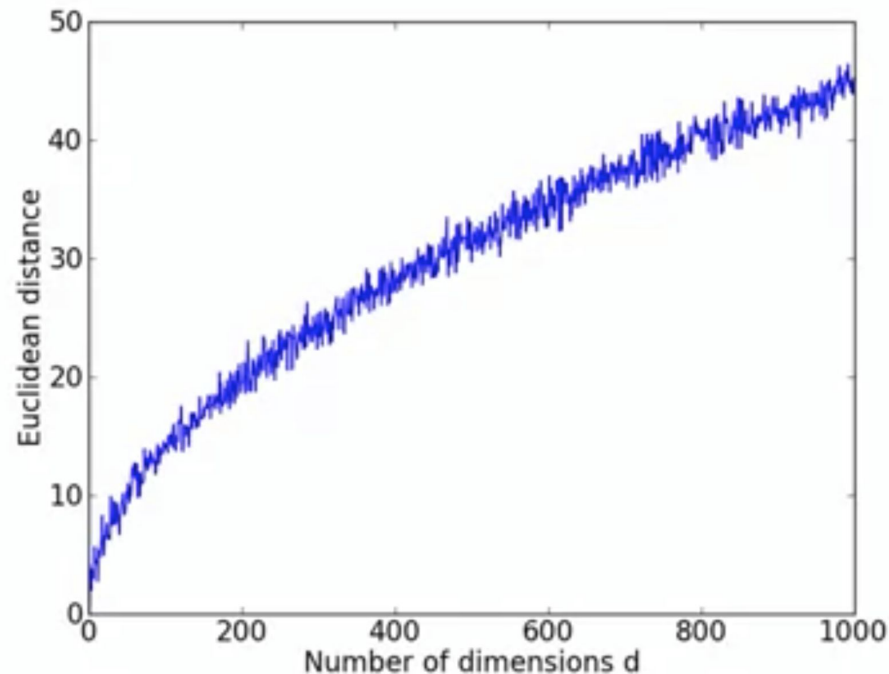
# Curse of dimensionality

- Increasing the number of features will not always improve classification accuracy.
- In practice, the inclusion of more features might actually lead to **worse** performance.
- When dimensionality increases, the volume of the space increases exponentially so that the data becomes sparse.
- The number of training examples required increases **exponentially** with dimensionality **d**.



# Curse of dimensionality (KNN example)

- In KNN, points tend to never be close together in high dimensional space



Credit: Bill Howe

# Curse of dimensionality (KNN example)

- Consider a **hypersphere** with radius  $r$  and dimension  $d$

- Distance between center and the corners is  $r\sqrt{d}$


- Consider **Square** with edge of length  $2r$

- $\sqrt{r^2 + r^2}$

- Consider **hypercube** with edge of length  $2r$


$$\frac{V_{\text{hypersphere}}}{V_{\text{hypercube}}} = \frac{\pi^{d/2}}{d2^{d-1}\Gamma(d/2)} \rightarrow 0 \text{ as } d \rightarrow \infty.$$

- Hypercube consist almost entirely of the “corners”



# Data Compression

- Reduces the required time and storage space
- Removing multi-collinearity improves the interpretation of the parameters of the machine learning model.



$$x^{(1)} \in R^2 \rightarrow z^{(1)} \in R$$

$$x^{(2)} \in R^2 \rightarrow z^{(1)} \in R$$

$\vdots$



$$x^{(m)} \in R^2 \rightarrow z^{(m)} \in R$$




# Data Compression

- Reduce data from 3D to 2D


In reality: 1000D to 100D



$x_2$




# Principal Component Analysis (PCA)



- Find a single line onto which to project this data

# Principal Component Analysis (PCA)





- Find a single line onto which to project this data

**minimize the projection error**

# Principal Component Analysis (PCA)


- PCA tries to find the surface (a straight line in this case) which has the minimum projection error



- Reduce n-D to k-D: find  $u^{(1)}, u^{(2)}, \dots, u^{(k)} \in R^n$  onto which to project the data, so as to minimize the projection error


# Feature scaling is important before PCA

- Normalization/Standardization is Important before PCA



# Linear regression vs. PCA

- Bonus point: What is the difference?



# Data pre-processing

- Training set:  $x^{(1)}, x^{(2)}, \dots, x^{(m)}$
- Preprocessing (feature scaling/mean normalization)

$$\mu_j = \frac{1}{m} \sum_i x_j^{(i)}$$

Replace each  $x_j^{(i)}$  with  $x_j - \mu_j$

If different features on different scales, scale features to have comparable range of values

$$x_j^{(i)} \leftarrow \frac{x_j^{(i)} - \mu_j}{s_j}$$

$s_j$ : Biggest – smallest, Standard deviation (more commonly)

# Principal Component Analysis Algorithm

- Goal: Reduce data from n-dimensions to k-dimensions
- Step 1: Compute “covariance matrix”

$$\Sigma = \frac{1}{m} \sum_{i=1}^n (x^{(i)})(x^{(i)})^T \quad \longrightarrow \quad \text{an } [n \times n] \text{ matrix}$$

- Step 2: Compute “eigenvectors” of the covariance matrix (e.g., using singular value decomposition)

$$[U, S, V] = \text{svd}(\Sigma)$$

$$\Sigma u_i = \lambda_i u_i$$

we assume eigenvalue  $\lambda_1 > \lambda_2 > \dots > \lambda_n$   
and  $u_1, u_2, \dots, u_n$  are the corresponding eigenvectors



$$U = \begin{bmatrix} | & | & & | \\ u^{(1)} & u^{(2)} & \dots & u^{(n)} \\ | & | & & | \end{bmatrix} \in \mathbb{R}^{n \times n}$$



# Principal Component Analysis Algorithm

- Goal: Reduce data from  $n$ -dimensions to  $k$ -dimensions
- **Principal components:**  $u^{(1)}, u^{(2)}, \dots, u^{(k)} \in \mathbb{R}^n$ 
  - Just take the first  $k$ -vectors from  $U$  (first  $k$  columns)

$$U = \begin{bmatrix} | & | & & | \\ u^{(1)} & u^{(2)} & \dots & u^{(n)} \\ | & | & & | \end{bmatrix} \in \mathbb{R}^{n \times n}$$

- $U_{reduce} = U(:, 1:k)$   an  $[n \times k]$  matrix
- $z = U_{reduce}^T * x$   a  $[k \times 1]$  vector

$$z = U_{reduce}^T x$$

# PCA algorithm summary

- Preprocessing
- Calculate sigma (covariance matrix)
- Calculate eigenvectors with **svd**
- Take k vectors from U ( $U_{reduce} = U(:, 1:k)$ )
- Calculate z ( $z = U_{reduce}^T x$ )

# Clustering


ESL 14.3

# Clustering Examples


- **Market segmentation**
  - group customers into different market segments
- **Organizing computer clusters** and data centers
  - for network layout and location
- **Astronomical data analysis**
  - understanding galaxy formation

# K-means


# Take unlabeled data and group into two clusters



# Randomly allocate $K$ points as cluster centroids




# Assign each point depending on which centroid it's closest to








# Move centroid to the average centroid of all points



# Iterate



# Animations



# K-means algorithm

- Input:
  - $K$  (number of clusters)
  - Training set  $\{x^{(1)}, x^{(2)}, x^{(3)}, \dots, x^{(m)}\}$

# K-means algorithm

- Randomly initialize  $K$  cluster centroids  $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^n$

Repeat{

for  $i = 1$  to  $m$

$c^{(i)} :=$  index (from 1 to  $K$ ) of cluster centroid  
closest to  $x^{(i)}$

Cluster assignment step

for  $k = 1$  to  $K$

$\mu_k :=$  average (mean) of points assigned to  
cluster  $k$

Centroid update step

}

# K-means optimization objective

- $c^{(i)}$  = Index of cluster (1, 2, ... K) to which example  $x^{(i)}$  is currently assigned
- $\mu_k$  = cluster centroid  $k$  ( $\mu_k \in \mathbb{R}^n$ )
- $\mu_{c^{(i)}}$  = cluster centroid of cluster to which example  $x^{(i)}$  has been assigned
- Optimization objective:

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

$$\min_{\substack{c^{(1)}, \dots, c^{(m)} \\ \mu_1, \dots, \mu_K}} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

**Example:**

$$x^{(i)} = 5$$

$$c^{(i)} = 5$$

$$\mu_{c^{(i)}} = \mu_5$$

# K-means algorithm

Randomly initialize  $K$  cluster centroids  $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^n$

Repeat{

for  $i = 1$  to  $m$

$c^{(i)} :=$  index (from 1 to  $K$ ) of cluster centroid  
closest to  $x^{(i)}$

**Cluster assignment step**

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

for  $k = 1$  to  $K$

cluster  $k$   $\mu_k :=$  average (mean) of points assigned to

**Centroid update step**

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

}


# K-means algorithm

- Stopping criterion
  - no (or minimum) re-assignments of data points to different clusters
  - no (or minimum) change of centroids
  - minimum decrease in the optimization objective
  - Maximum iterations




# Random initialization

- Randomly pick  $K$  training examples
- Set  $\mu_1, \mu_2, \dots, \mu_K$  equal to those  $K$  examples



# Local Optimum



# Solutions to Local Optimum

- Multiple random initialization
- Furthest Point Heuristic

# 1) Multiple random initialization

For  $i = 1$  to 100 {

Randomly initialize K-means.

Run K-means. Get  $c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K$

Compute the cost function (distortion)

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

}

Pick clustering that gave the lowest cost

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

## 2) Furthest Point Heuristic

Choose  $\mu_1$  arbitrarily (or at random)


For  $j = 2$  to  $K$

Pick  $\mu_j$  among data points  $x^{(1)}, x^{(2)}, \dots, x^{(m)}$  that is farthest from previously chosen  $\mu_1, \mu_2, \dots, \mu_{j-1}$

Slide credit: Maria-Florina Balcan

# How to choose K?


- Cluster validity indices!!
  - An assessment of how “good” a given clustering is
  - Generally based on inter-cluster separation and intra-cluster spread
  - Ex: Davies-Bouldin Index, Silhouette Score, etc.
- Can also use visualizations to do it manually



Credit: Andrew Ng

# How to choose K?


- 1) Elbow method



Credit: Andrew Ng

# How to choose K?

- 2) Try multiple K and use performance from downstream tasks for selection (ie. “engineering decision”)



**K=3 or K=5?**

- cost of making extra sizes vs. how well distributed the products are
- How important are those sizes though?

Credit: Andrew Ng



# Summary

- k-means is the most popular clustering algorithm
- Pros
  - Simple: easy to understand and to implement
  - Efficient
    - Time complexity:  $O(tkn)$ , where  $n$  is the number of data points,  $k$  is the number of clusters, and  $t$  is the number of iterations.
    - Since both  $k$  and  $t$  are small. k-means is considered a linear algorithm.
- Cons
  - The user needs to specify  $k$
  - Sensitive to outliers
  - Not suitable for special data structure



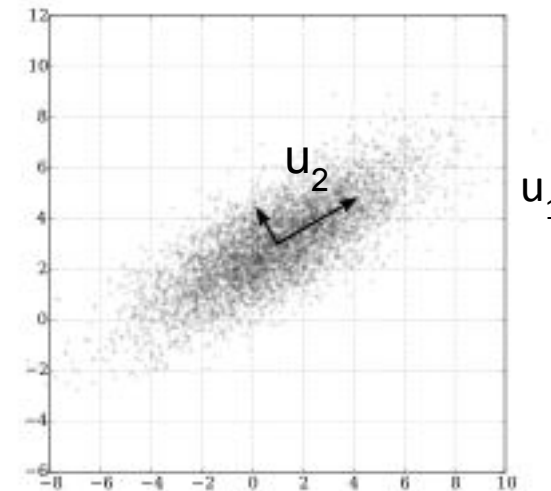
**Michigan Tech**

# Geometric interpretation of PCA

- PCA chooses the **eigenvectors** of the covariance matrix corresponding to the **largest** eigenvalues.
- The **eigenvalues** correspond to the **variance** of the data along the eigenvector directions.
- Therefore, PCA projects the data along the directions where the data varies **most**.


$u_1$ : direction of max variance

$u_2$ : orthogonal to  $u_1$



# Reconstruction from compressed representation

- Compression:  $z^{(i)} = U_{\text{reduce}}^T x^{(i)}$
- Reconstruction:  $x_{\text{approx}}^{(i)} = U_{\text{reduce}} z^{(i)}$
- $x_{\text{approx}}^{(i)} \in R^n$   $U_{\text{reduce}} \in R^{n \times k}$



# Application of PCA

- Compression
  - Reduce memory/disk needed to store data
  - Speed up learning algorithm
- Visualization ( $k=2$ ,  $k=3$ )
- Bad use of PCA
  - Use it to prevent over-fitting
    - Fewer features ☐ less likely overfitting?
    - Might work OK, but not a good way to address overfitting
    - PCA does not look at  $y$ , may throw away some critical information
    - Use regularization instead.



**Michigan Tech**

# Questions + Comments?

