

CS5841/EE5841 Machine Learning

Lecture 12a:NN notes

Evan Lucas



Michigan Tech

Loss functions

- Cross Entropy
 - Includes softmax(), just use linear output layer
 - For multi-class classification
- Binary Cross Entropy
 - For single output neuron binary classification
- Mean Squared Error
 - Regression, for single output neuron
- Others



Wait, when do we use softmax?

- Don't include it in your network because it's already included in pytorch's cross entropy loss definition
- Inference
 - If you just want top predicted class - it doesn't matter
 - If you want top-k or prediction probabilities, compute outputs and run through



Why shouldn't we include softmax()?

- Numerical stability
 - underflow - really small numbers get approximated as zero
 - overflow - really big numbers get approximated as infinity



What comes out of a neural network?

- Logits
 - Raw, unnormalized values
- If multi-class classification, these can be treated as independent classifier outputs
 - This is why softmax is helpful - normalizing w.r.t. other class predictions
- After normalization - prediction probability distribution
- What is the probability distribution of labels?



PyTorch Datasets and Dataloaders

- use `next(iter(myDataloader))` to see the output of the dataloader for debugging purposes
- PyTorch has some good tutorials for building custom Datasets as subclass of PyTorch Dataset
 - There's also a lazy way...
-



Lazy dataset creation (won't always work)

```
import torch.utils.data as data_utils
```

```
train = data_utils.TensorDataset(features, targets)
```

```
train_loader = data_utils.DataLoader(train,  
    batch_size=50, shuffle=True)
```



Easy test/train split in pytorch

```
train_set, val_set =  
    torch.utils.data.random_split(dataset, [train_size,  
    test_size])
```

Other ways to specify split



Questions + Comments?



Resources used

http://cs231n.stanford.edu/slides/2023/lecture_2.pdf

