

Homework 2 - KNN and Logistic Regression

EE/CS5841, Spring 2024

```
1 # References:
2 # I have used many references for this assisgment
3 # 1.
  https://web.stanford.edu/class/archive/cs/cs109/cs109.1166/pdfs/40%20LogisticRegression.pdf
4 # 2. Cs229 lectures
5 # 3. https://cseweb.ucsd.edu/~elkan/250B/logreg.pdf
6 # 4. https://zlatankr.github.io/posts/2017/03/06/mle-gradient-descent
7 # 5. https://zstevenwu.com/courses/s20/csci5525/resources/slides/lecture05.pdf
8 # 6. https://fluxml.ai/Flux.jl/stable/tutorials/logistic_regression/
9 # 7. https://machinelearninggeek.com/mnist-with-julia/
10 # 8. https://int8.io/category/classification/
11 # 9. https://int8.io/logstic-regression-with-gradient-descent-in-julia/
12 # 10. https://int8.io/logistic-regression-part-ii-evaluation/
13 # 11. I have various Large language models for code understanding and generation
    julia and of debugging too
```

Template to load MNIST

Below is an example template in Julia for loading the MNIST dataset and plotting a couple of images. (Make sure that you have the necessary packages installed, and adjust the paths accordingly)

```
1 begin
2     using Plots
3     using MLDatasets
4     using LinearAlgebra
5 end
```


Selection deleted

(28, 28, 60000)

```

1 begin
2   # load training set
3   train_x, train_y = MNIST(split=:train)[: ]
4   train_x = permutedims(train_x, (2,1,3))
5
6   # load test set
7   test_x, test_y = MNIST(split=:test)[: ]
8   test_x = permutedims(test_x, (2,1,3))
9   size(train_x)
10 end

```



```

1 begin
2   # Plot a couple of training images. Note that the image y axis is flipped, this
   is common for images because of the way pixel position is notated.
3   plot(heatmap(train_x[:, :, 1], c=:binary, yflip=true), heatmap(train_x[:, :, 2],
   c=:binary, yflip=true), heatmap(train_x[:, :, 3], c=:binary,
   yflip=true), heatmap(train_x[:, :, 4], c=:binary, yflip=true))
4
5 end

```

Selection deleted
[5, 0, 4, 1]

```
1 train_y[1:4]
```

(10 pts) Part 1: 1-nearest neighbour classifier

Use the MNIST dataset of handwritten digits with corresponding labels. It may be helpful to convert the 28x28 pixel image into a 784-dimensional feature vector. Implement a 1-nearest neighbor classifier without scaling or normalizing pixel values. The goal is to identify the number of images in the testing set that have been correctly labelled. Discuss the results and any insights gained from the analysis.

one_nearest_neighbor_classifier (generic function with 1 method)

```

1 begin
2     function one_nearest_neighbor_classifier(train_x, train_y, test_x)
3         # Your code for 1-nearest neighbor classifier here
4         # Loop over testing_data, find nearest neighbor in training_data, and
          check correctness
5         infer_labels=[]
6         for i in 1:size(test_x, 3)
7
8             # Flatten the current test image into a 784-dimensional vector
9             test_image = reshape(test_x[:, :, i], 784)
10
11            # Find the index of the training image with the minimum Euclidean
          distance
12            closest_index = argmin([LinearAlgebra.norm(test_image -
          reshape(train_x[:, :, j], 784)) for j in 1:size(train_x, 3)])
13
14            # Infer the label from the training set and store it
15            push!(infer_labels, train_y[closest_index])
16        end
17        return infer_labels
18    end
19 end

```

accuracy (generic function with 1 method)

```

1 begin
2     function accuracy(infer_labels, true_labels)
3         # Your code here
4         accuracy = sum(infer_labels .== true_labels) / length(true_labels)
5         println(accuracy)
6         return accuracy
7     end
8 end

```

0.9691

```
1 accuracy(one_nearest_neighbor_classifier(train_x, train_y, test_x), test_y)  
2
```

0.9691



Selection deleted

(10 pts) Part 2: KNN leave-one-out

Implement a KNN leave-one-out approach and test values of K from 1 to 20. Plot the leave-one-out error vs. K. Add a comment discussing the results. (If you are running into time problems using all 60,000 data points for leave-one-out, feel free to randomly sample the training set to estimate the best K.)

Selection deleted


```

knn_leave_one_out (generic function with 1 method)

1 begin
2 using Random
3 using Statistics
4 using StatsBase
5     # Set a fixed seed for reproducibility
6     seed_value = 42
7     Random.seed!(seed_value)
8
9     # Define your sample size
10    sample_size = 100
11
12    # Randomly sample indices for both training and test sets
13    train_indices = randperm(size(train_x, 3))[1:sample_size]
14    test_indices = randperm(size(test_x, 3))[1:sample_size]
15
16    # Use the sampled indices to extract the corresponding data
17    sampled_train_x = train_x[:, :, train_indices]
18    sampled_train_y = train_y[train_indices]
19
20    sampled_test_x = test_x[:, :, test_indices]
21    sampled_test_y = test_y[test_indices]
22
23    function knn_leave_one_out(train_x, train_y, k_values)
24        # Your code for KNN leave-one-out approach here
25        infer_labels = []
26
27        for k in k_values
28            inferred_labels = []
29            correct_predictions = 0
30
31            for i in 1:sample_size
32                # Flatten the current test image into a 784-dimensional vector
33                test_image = reshape(sampled_test_x[:, :, i], 784)
34
35                # Find the indices and labels of the k-nearest neighbors
36                distances = [LinearAlgebra.norm(test_image - reshape(train_x[:, :,
37j], 784)) for j in 1:sample_size]
38                nearest_indices = sortperm(distances)[1:k]
39                nearest_labels = train_y[nearest_indices]
40
41                # Infer the label by majority voting
42                inferred_label = argmax(countmap(nearest_labels))
43                push!(inferred_labels, inferred_label)
44
45                push!(infer_labels, inferred_labels)
46            end
47            return infer_labels
48        end
49    end
50 end

```

WARNING: using StatsBase.crossentropy in module workspace#4 conflicts with an existing identifier.



```
1 begin
2   #Plot your accuracy vs k here
3   k = 1:20
4   accur=[]
5   for i in k
6     predicted_y=knn_leave_one_out(sampled_train_x,sampled_train_y,k)
7     acc = sum(predicted_y[i].==sampled_test_y)/length(sampled_test_y)
8     push!(accur, acc)
9   end
10   #replace this with your real accuracy!!
11   plot(k,accur,xlabel="k",ylabel="Accuracy",)
12
13 end
14
15
```

1 Enter cell code...

Selection deleted

(30 pts) Part 3: Downsampling and KNN Experiment

(10 pts) Section a: Downsampling

Implement a downsampling function for an image dataset and conducting a K-nearest neighbors (KNN) leave-one-out experiment. The goal is to observe the impact of downsampling on the testing results and subjectively evaluate the query time of the classifier.

The function `downsample_image` takes an image represented as a 784-dimensional feature vector (1D array of UInt8) and a downsampling factor `n`. The function should downsample the image by selecting every `n`th pixel (feature) and return the downsampled feature vector.

`downsample_image` (generic function with 1 method)

```
1 function downsample_image(image, n)
2     # Your code for downsampling here
3     # Downsample by selecting every nth pixel along both dimensions
4     downsampled_image = image[1:n:end, 1:n:end]
5
6     # Reshape the downsampled image back into a 1D array
7     return vec(downsampled_image)
8 end
9
```

Selection deleted

knn_leave_one_out_experiment (generic function with 1 method)

```
1 begin
2     function knn_leave_one_out_experiment(train_x, train_y, test_x, test_y, k_values)
3         # Your code for the KNN leave-one-out experiment here
4         downsample_factor= 10
5         accuracy_vector = []
6
7         for k in k_values
8             inferred_labels = []
9
10            for i in 1:size(test_x, 3)
11                # Flatten and downsample the current test image
12                test_image = downsample_image(reshape(sampled_test_x[:, :, i], 784),
downsample_factor)
13
14                # Downsample training images as well
15                sampled_train_x_downsampled =
[downsample_image(reshape(sampled_train_x[:, :, j], 784), downsample_factor) for j
in 1:size(train_x, 3)]
16
17                # Find the indices and labels of the k-nearest neighbors
18                distances = [LinearAlgebra.norm(test_image -
sampled_train_x_downsampled[j]) for j in 1:size(train_x, 3)]
19                nearest_indices = sortperm(distances)[1:k]
20                nearest_labels = sampled_train_y[nearest_indices]
21
22                # Infer the label by majority voting
23                inferred_label = argmax(countmap((nearest_labels)))
24                push!(inferred_labels, inferred_label)
25            end
26
27            # Calculate accuracy and store in accuracy_vector
28            acc = accuracy(inferred_labels, sampled_test_y)
29            push!(accuracy_vector, acc)
30        end
31
32        return accuracy_vector
33    end
34 end
35
```

Selection deleted

```
[0.68, 0.63, 0.62, 0.62, 0.59, 0.59, 0.58, 0.58, 0.58, 0.58, 0.56, 0.56, 0.52, 0.49, 0.5, 0
```

```
1 knn_leave_one_out_experiment(sampled_train_x, sampled_train_y, sampled_test_x,  
    sampled_test_y, k)
```

```
0.68  
0.63  
0.62  
0.62  
0.59  
0.59  
0.58  
0.58  
0.58  
0.58  
0.56  
0.56  
0.52  
0.49  
0.5  
0.51  
0.48  
0.47  
0.47  
0.47
```



(20 pts) section b: Bin-based Downsampling

Implement a function that downsamples the image using cluster-based downsampling. For example, if n is 4, bin the pixels into groups of 4 and replace each pixel in a cluster with the average value of the pixels in that cluster. Create a second downsampling function that takes the the max value of the pixels in the bin. Comment on the testing results and the query time of the classifier. Start from the 28x28 images, bin them (maybe try displaying what they look like downsampled if you're curious), and then reshape them into a 1xm vector before using your classifier.

Try a bin size of 2x2 and 4x4 for full credit. If you're really curious, feel free to try other sizes.

Selection deleted

average_pooled_downsampling (generic function with 1 method)

```
1 begin
2     function average_pooled_downsampling(image, n)
3         # Your code for clustering and downsampling here
4         height, width, num_samples = size(image)
5
6         # Calculate the number of clusters in each dimension
7         clusters_per_row = div(width, n)
8         clusters_per_col = div(height, n)
9
10        # Initialize the downsampled image
11        downsampled_image = zeros(Float64, clusters_per_col, clusters_per_row,
num_samples)
12
13        for k in 1:num_samples
14            for i in 1:clusters_per_col
15                for j in 1:clusters_per_row
16                    # Define the boundaries of the current cluster
17                    start_row, end_row = (i - 1) * n + 1, i * n
18                    start_col, end_col = (j - 1) * n + 1, j * n
19
20                    # Extract the pixels within the current cluster
21                    cluster_pixels = image[start_row:end_row, start_col:end_col, k]
22
23                    # Calculate the average value for the cluster
24                    average_value = mean(cluster_pixels)
25
26                    # Assign the average value to the downsampled image
27                    downsampled_image[i, j, k] = average_value
28                end
29            end
30        end
31
32        return downsampled_image
33    end
34 end
```

Selection deleted

```

0.0242647 0.101716 0.0 0.257353 0.823284 0.0 0.0
0.0 0.0 0.0 0.457108 0.497059 0.0 0.0
0.0 0.0 0.0 0.678922 0.456373 0.0 0.0
0.0 0.0 0.0 0.339706 0.104902 0.0 0.0

```

```

;;; ...

```

```

[:, :, 98] =
0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0419118 0.715686 0.376961 0.0 0.0
0.0 0.0 0.41299 0.118382 0.575735 0.0 0.0
0.0 0.0 0.328676 0.707598 0.443627 0.0 0.0
0.0 0.0 0.527451 0.0404412 0.407108 0.0 0.0
0.0 0.0 0.520343 0.400245 0.334804 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0

```

```

[:, :, 99] =
0.0 0.0 0.0669118 0.201471 0.0 0.0 0.0
0.0 0.0232843 0.693382 0.669608 0.121324 0.0 0.0
0.0 0.0 0.0193627 0.409559 0.255637 0.0 0.0
0.0 0.0 0.0 0.66299 0.0335784 0.00784314 0.0
0.0 0.0 0.36348 0.931128 0.899755 0.478186 0.0
0.0 0.0 0.520343 0.393873 0.0958333 0.106373 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0

```

```

[:, :, 100] =
0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.00833333 0.332843 0.542157 0.332353 0.00833333 0.0
0.0 0.640931 0.555882 0.25098 0.779902 0.707353 0.0
0.0 0.902206 0.0 0.0 0.0 0.902941 0.0
0.0 0.680637 0.557353 0.361765 0.626226 0.683333 0.0
0.0 0.0720588 0.476961 0.540196 0.353676 0.00808824 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0

```

```

1 average_pooled_downsampling(sampled_train_x, 4)

```

Selection deleted

14x14x100 Array{Float64, 3}:

```
[ :, :, 1] =
0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0284314 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.896078 0.223529 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.992157 0.859804 0.027451 0.0 0.0
0.0 0.0 0.0 0.0 0.626471 0.988235 0.272549 0.0 0.0
0.0 0.0 0.0 0.0 ... 0.497059 0.990196 0.434314 0.0 0.0
0.0 0.0 0.0656863 0.558824 0.361765 0.988235 0.494118 0.0 0.0
0.0 0.00392157 0.654902 0.992157 0.42451 0.990196 0.495098 0.0 0.0
0.0 0.227451 0.988235 0.67451 0.591177 0.988235 0.407843 0.0 0.0
0.0 0.498039 0.990196 0.965686 0.992157 0.857843 0.027451 0.0 0.0
0.0 0.095098 0.605882 0.907843 ... 0.714706 0.222549 0.0 0.0 0.0
0.0 0.0 0.0 0.027451 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
```

```
[ :, :, 2] =
0.0 0.0 0.0 0.0 0.0 ... 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.287255 0.0156863 0.0 0.0 0.0
0.0 0.0 0.0 0.132353 0.914706 ... 0.221569 0.0 0.0 0.0
0.0 0.0 0.0 0.00588235 0.0490196 0.288235 0.0 0.0 0.0
0.0 0.0 0.0 0.177451 0.669608 0.0892157 0.0 0.110784 0.0
0.0 0.0 0.0 0.876471 0.293137 0.798039 0.737255 0.660784 0.0
0.0 0.0 0.0 0.671569 0.611765 0.0147059 0.0 0.0 0.0
0.0 0.0 0.0 0.0117647 0.147059 ... 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0 0.0
```

```
1 average_pooled_downsampling(sampled_train_x, 2)
```

Selection deleted

max_pooled_downsampling (generic function with 1 method)

```
1 begin
2 function max_pooled_downsampling(image, n)
3     # Your code for clustering and downsampling here
4     # Get the dimensions of the original image
5     height, width, num_samples = size(image)
6
7     # Calculate the number of clusters in each dimension
8     clusters_per_row = div(width, n)
9     clusters_per_col = div(height, n)
10
11     # Initialize the downsampled image
12     downsampled_image = zeros(Float64, clusters_per_col, clusters_per_row, num_samples)
13
14     # Iterate over clusters and calculate max value for each
15     for k in 1:num_samples
16         for i in 1:clusters_per_col
17             for j in 1:clusters_per_row
18                 # Define the boundaries of the current cluster
19                 start_row, end_row = (i - 1) * n + 1, i * n
20                 start_col, end_col = (j - 1) * n + 1, j * n
21
22                 # Extract the pixels within the current cluster
23                 cluster_pixels = image[start_row:end_row, start_col:end_col, k]
24
25                 # Calculate the max value for the cluster
26                 max_value = maximum(cluster_pixels)
27
28                 # Assign the max value to the downsampled image
29                 downsampled_image[i, j, k] = max_value
30             end
31         end
32     end
33
34     return downsampled_image
35 end
36 end
```

Selection deleted

```
7x7x100 Array{Float64, 3}:
```

```
[:, :, 1] =
 0.0      0.0      0.0      0.0      0.113725  0.0      0.0
 0.0      0.0      0.0      0.996078  0.996078  0.992157  0.0
 0.0      0.0      0.996078  0.992157  0.996078  0.992157  0.0
 0.0156863 0.996078  0.992157  0.941176  0.94902   0.992157  0.0
 1.0      0.996078  0.996078  0.996078  0.996078  0.992157  0.0
 0.380392 0.992157  0.992157  0.992157  0.992157  0.741176  0.0
 0.0      0.0      0.0      0.0      0.0      0.0      0.0
```

```
[:, :, 2] =
 0.0  0.0      0.0      0.0      0.0      0.0      0.0
 0.0  0.0      0.0      0.298039  0.0      0.0      0.0
 0.0  0.423529  0.996078  0.996078  0.996078  0.443137  0.0
 0.0  0.670588  0.996078  0.917647  1.0      0.619608  0.329412
 0.0  1.0      0.996078  0.996078  0.996078  0.941176  0.898039
 0.0  0.0470588 0.294118  0.0      0.0      0.0      0.0
 0.0  0.0      0.0      0.0      0.0      0.0      0.0
```

```
[:, :, 3] =
 0.0      0.0      0.0      0.0      0.0      0.0      0.0
 0.0      0.0      0.121569  0.992157  1.0      0.972549  0.0
 0.913725 0.996078  0.996078  0.996078  0.996078  0.972549  0.0
 0.388235 0.662745  0.0      0.996078  1.0      0.0      0.0
 0.0      0.0      0.0      0.996078  0.996078  0.0      0.0
 0.0      0.0      0.0      1.0      0.996078  0.0      0.0
 0.0      0.0      0.0      1.0      0.662745  0.0      0.0
```

```
;;; ...
```

```
[:, :, 98] =
 0.0  0.0  0.0      0.0      0.0      0.0  0.0
 0.0  0.0  0.576471  1.0      0.992157  0.0  0.0
 0.0  0.0  0.992157  0.996078  0.996078  0.0  0.0
```

```
1 max_pooled_downsampling(sampled_train_x, 4)
```

Selection deleted

(40 pts) Part 4: Regularized Logistic regression classifier

Build a regularized logistic regression classifier, where you use ridge (ℓ_2) regularization. Test this classifier on the MNIST data set by developing 10 classifiers: 0 versus all, 1 versus all, 2 versus all, ... , 9 versus all. Provide a confusion matrix, accuracy for each digit, and overall accuracy. Plot the overall test accuracy versus the regularization value where a log-scale is used for regularization value.

Essentially, the '1 versus all' classifier is trained to give you a probability of the digit 1 versus all other digits. Hence, digit 1 is class +1 and all other digits are class 0. Hence, to classify a test image, you take the maximum probability from all 10 classifiers, giving the predicted class of the input image. ℓ_2 regularized logistic regression uses the following log-likelihood,

$$L(\mathbf{w}) = \sum_{i=1}^N \log(1 + \exp(-\mathbf{y}_i \mathbf{w}^T \mathbf{x}_i)) + \lambda \|\mathbf{w}\|_2$$

Selection deleted

$$\mathcal{L}(\mathbf{w}) = \sum_{i=1}^N \log(1 + \exp(-y_i \mathbf{w}^T \mathbf{x}_i)) + \frac{\lambda}{2} \|\mathbf{w}\|_2.$$

```

1
2 md"""
3 # (40 pts) Part 4: Regularized Logistic regression classifier
4 Build a regularized logistic regression classifier, where you use ridge ('2)
5 regularization. Test this classifier
6 on the MNIST data set by developing 10 classifiers: 0 versus all, 1 versus all, 2
7 versus all, ... , 9 versus all. Provide a confusion
8 matrix, accuracy for each digit, and overall accuracy. Plot the overall test
9 accuracy versus the regularization value where a
10 log-scale is used for regularization value.
11
12 Essentially, the '1 versus all' classifier is trained to give you a probability
13 of the digit 1 versus all other digits. Hence, digit
14 1 is class +1 and all other digits are class 0. Hence, to classify a test image,
15 you take the maximum probability from all 10
16 classifiers, giving the predicted class of the input image.
17 '2 regularized logistic regression uses the following log-likelihood,
18
19 >### L(w) = \sum_{i=1}^N \log(1+\exp(-y_i w^T x_i)) + \lambda \|w\|_2
20 #
21 ![eq]
22 (https://miro.medium.com/v2/resize:fit:640/format:webp/1\*lmBn5fQTXB4gxqFylmhGqQ.png)
23
24 """
25

```

```

1 begin
2   using Flux: Chain, Dense, σ, binarycrossentropy, gradient, Optimise, update!,
3   params, onehotbatch, onecold, crossentropy, normalise
4   using Flux
5   using Printf
6   using LinearAlgebra: norm
7 end

```

```

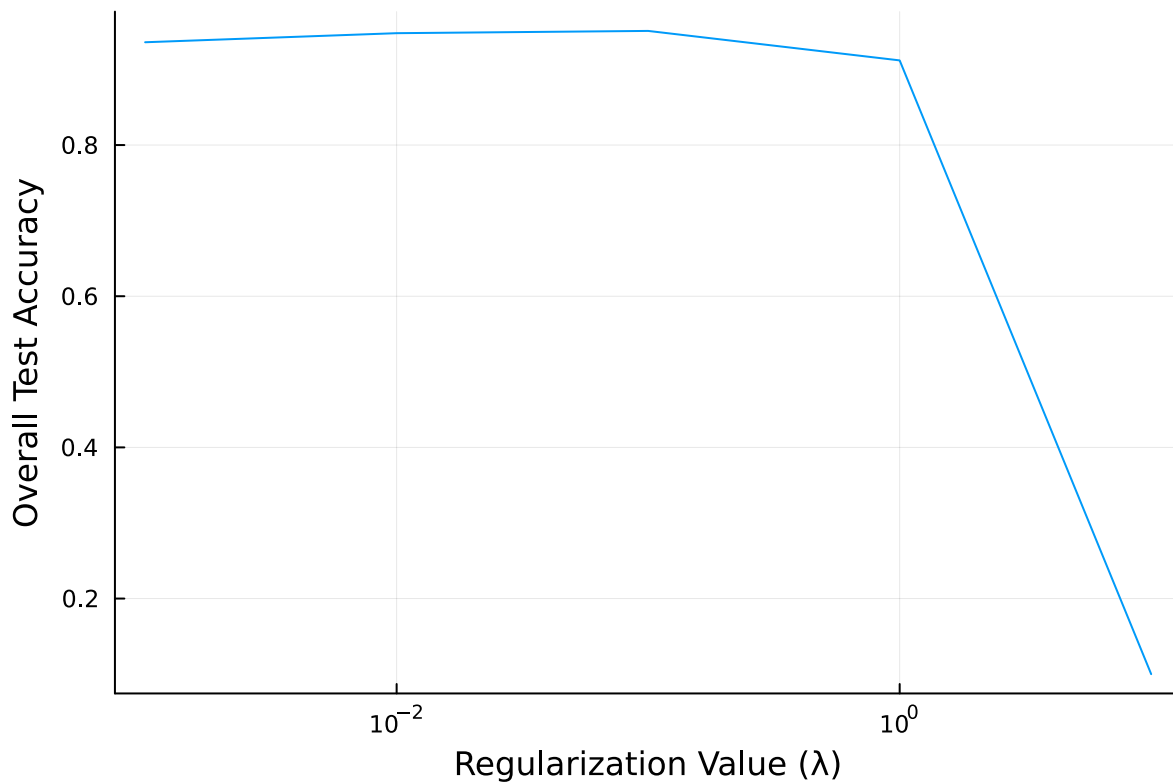
1 begin
2 # Preprocess data (normalize each image individually)
3 function preprocess_data(images, labels, digit)
4     X = Float32.(images) # Convert to Float32
5     X = reshape(X, :, size(X, 3))
6     X = [(x .- mean(x)) / std(x) for x in eachcol(X)] # Normalize each image
        individually
7     X = hcat(X...)
8     y = (labels .== digit)
9     return X, y
10 end
11
12 # Function to calculate confusion matrix
13 function calculate_confusion_matrix(predictions, actual)
14     num_classes = 2
15     confusion_matrix = zeros{Int, num_classes, num_classes}
16
17     for i in 1:length(actual)
18         true_class = Int(actual[i]) + 1
19         pred_class = Int(predictions[i]) + 1
20         confusion_matrix[true_class, pred_class] += 1
21     end
22
23     return confusion_matrix
24 end
25
26 # Define logistic regression model
27 function logistic_regression_model(input_size)
28     return Chain(Dense(input_size, 1, σ))
29 end
30
31 # Define loss function (logistic loss + ridge regularization)
32 function loss_function(model, X, y, λ)
33     ŷ = model(X)
34     ŷ = vec(ŷ) # Reshape ŷ to match the shape of y
35     loss = Flux.binarycrossentropy(ŷ, y)
36     reg = sum(norm(param, 2)^2 for param in params(model))
37     return loss + λ * reg
38 end
39
40 # Define training function using gradient descent
41 function train_model!(model, X, y, λ, lr, epochs)
42     opt = Flux.Optimise.Descent(lr)
43     ps = params(model)
44     dataset = Iterators.repeated((X, y), epochs)
45     for (X_batch, y_batch) in dataset
46         grads = Flux.gradient(() -> loss_function(model, X_batch, y_batch, λ), ps)
47         Flux.Optimise.update!(opt, ps, grads)
48     end
49 end
50
51 # Train classifiers for each digit
52 function train_classifiers(train_x, train_y, test_x, test_y, λ, lr, epochs)
53     accuracies = []

```

```
54  confusion_matrices = []
55  for digit in 0:9 # Binary classification
56      println("Training classifier for digit ", digit)
57      X_train, y_train = preprocess_data(train_x, train_y, digit)
58      X_test, y_test = preprocess_data(test_x, test_y, digit)
59      model = logistic_regression_model(size(X_train, 1))
60      train_model!(model, X_train, y_train, λ, lr, epochs)
61
62      # Check if there are any weight parameters
63      weights = params(model)
64      if isempty(weights)
65          println("No weight parameters found for classifier $digit")
66          push!(accuracies, NaN)
67          continue
68      end
69
70      ŷ = model(X_test)
71      predictions = Flux.round.(Int, vec(ŷ))
72      accuracy = mean(predictions .== y_test)
73      push!(accuracies, accuracy)
74
75      # Calculate and store confusion matrix
76      confusion_matrix = calculate_confusion_matrix(predictions, y_test)
77      push!(confusion_matrices, confusion_matrix)
78
79      println("Accuracy for digit ", digit, ": ", accuracy)
80      println("Confusion Matrix for digit ", digit, ":")
81      println(confusion_matrix)
82  end
83  return accuracies, confusion_matrices
84 end
85
86 # Define parameters
87 samp_size = 60000
88 λ_vals = [0.001, 0.01, 0.1, 1.0, 10]
89 learning_rate = 0.1
90 epochs = 10
91
92 # Train classifiers
93 accuracies_per_lambda = []
94 confusion_matrices_per_lambda = []
95
96 for λ in λ_vals
97     println("Lambda = ", λ)
98     accuracies, confusion_matrices = train_classifiers(sampled_train_x,
99 sampled_train_y, sampled_test_x, sampled_test_y, λ, learning_rate, epochs)
100     push!(accuracies_per_lambda, accuracies)
101     push!(confusion_matrices_per_lambda, confusion_matrices)
102 end
103
104 # Display overall accuracy for all lambdas
105 println("Overall Accuracies for All Lambdas:")
106 for (λ, accs) in zip(λ_vals, accuracies_per_lambda)
107     println("Lambda = ", λ, ": ", mean(accs))
108 end
```

107 end

```
Lambda = 0.001
Training classifier for digit 0
Accuracy for digit 0: 0.96
Confusion Matrix for digit 0:
[90 4; 0 6]
Training classifier for digit 1
Accuracy for digit 1: 0.98
Confusion Matrix for digit 1:
[87 1; 1 11]
Training classifier for digit 2
Accuracy for digit 2: 0.98
Confusion Matrix for digit 2:
[87 0; 2 11]
Training classifier for digit 3
Accuracy for digit 3: 0.88
Confusion Matrix for digit 3:
[85 2; 10 3]
Training classifier for digit 4
Accuracy for digit 4: 0.91
Confusion Matrix for digit 4:
[89 0; 9 2]
Training classifier for digit 5
Accuracy for digit 5: 0.9
Confusion Matrix for digit 5:
[85 3; 7 5]
Training classifier for digit 6
Accuracy for digit 6: 0.96
Confusion Matrix for digit 6:
[91 0; 4 5]
Training classifier for digit 7
Accuracy for digit 7: 0.95
Confusion Matrix for digit 7:
```



```

1 begin
2 # Plot overall test accuracy versus regularization value
3 plot( $\lambda$ _vals, mean.(accuracies_per_lambda), xscale=:log10, xlabel="Regularization
  Value ( $\lambda$ )", ylabel="Overall Test Accuracy", label="")
4
5 end

```

InterruptedException:

```

1 md"""
2 The general weight update equation for gradient ascent is the following:
3
4 >###  $w_{t+1} = w_t + \eta_t \nabla L(f(x;w), y)$ 
5 #
6
7 Refer to the uploaded file '40 Logistic Regression.pdf' for a derivation of the
  gradient of unregularized logistic regression
8
9
10
11
12 ### (30 pts) section a: Regularized logistic regression
13 Derive the update equation for the regularized logistic regression by taking the
  gradient of the maximum log likelihood and present your derivation in a few
  sentences. (10 points)
14
15 Apply that classifier to the
16 MNIST data set. Experiment with the learning rate ( $\eta_t$ ) to train your model. Don't
  forget to add a column of 1s to your 784-length feature vector (the image values) so
  that you get the
17 bias term with your classifier. (20 points)
18

```

```
19 Use the given test-train split with no cross validation.  
20  
21  
22  
23 """
```

InterruptException:

```
1 # I am uploading the derivation in another file please check it out
```

```
1 begin
2   using Flux: ADAM
3   using Random: seed!
4
5   # Preprocess data (normalize each image individually)
6   function new_preprocess_data(images, labels, digit)
7       X = Float32.(images) # Convert to Float32
8       X = reshape(X, :, size(X, 3))
9       X = [(x .- mean(x)) / std(x) for x in eachcol(X)] # Normalize each image
10      individually
11      X = hcat(X...)
12      X = vcat(X, ones(1, size(X, 2))) # Add a column of 1s for bias
13      y = (labels .== digit)
14      return X, y
15  end
16
17  # Define logistic regression model with bias
18  function new_logistic_regression_model(input_size)
19      return Chain(Dense(input_size, 1, σ))
20  end
21
22  # Define loss function (logistic loss + ridge regularization)
23  function new_loss_function(model, X, y, λ)
24      ŷ = model(X)
25      ŷ = vec(ŷ)
26      loss = Flux.binarycrossentropy(ŷ, y)
27      reg = sum(norm(param, 2)^2 for param in params(model))
28      return loss + λ * reg # New equation for regularization
29  end
30
31  # Define training function using gradient descent
32  function new_train_model!(model, X, y, λ, lr, epochs)
33      opt = ADAM(lr)
34      ps = params(model)
35      dataset = Iterators.repeated((X, y), epochs)
36      for (X_batch, y_batch) in dataset
37          grads = Flux.gradient(() -> new_loss_function(model, X_batch, y_batch,
38              λ), ps)
39          Flux.Optimise.update!(opt, ps, grads)
40      end
41  end
42
43  # Train classifiers for each digit with bias
44  function train_classifiers_with_bias(train_x, train_y, test_x, test_y, λ, lr,
45      epochs)
46      accuracies = []
47      confusion_matrices = []
48      for digit in 0:9 # Binary classification
49          println("Training classifier for digit ", digit)
50          X_train, y_train = new_preprocess_data(train_x, train_y, digit)
51          X_test, y_test = new_preprocess_data(test_x, test_y, digit)
52          model = new_logistic_regression_model(size(X_train, 1))
53          new_train_model!(model, X_train, y_train, λ, lr, epochs)
```

```
52
53     # Check if there are any weight parameters
54     weights = params(model)
55     if isempty(weights)
56         println("No weight parameters found for classifier $digit")
57         push!(accuracies, NaN)
58         continue
59     end
60
61
62     ŷ = model(X_test)
63     predictions = Flux.round.(Int, vec(ŷ))
64     accuracy = mean(predictions .== y_test)
65     push!(accuracies, accuracy)
66
67     # Calculate and store confusion matrix
68     confusion_matrix = calculate_confusion_matrix(predictions, y_test)
69     push!(confusion_matrices, confusion_matrix)
70
71     println("Accuracy for digit ", digit, ": ", accuracy)
72     println("Confusion Matrix for digit ", digit, ":")
73     println(confusion_matrix)
74 end
75 return accuracies, confusion_matrices
76 end
77
78 # Define parameters
79
80 λ_val = 0.001
81 learning_rates = [0.001, 0.01, 0.1, 1.0, 10]
82
83
84 # Train classifiers with bias
85 new_accuracies_per_learning_rate = []
86 new_weights_per_learning_rate = []
87 for learning_rate in learning_rates
88     println("learning_rate = ", learning_rate)
89     accuracies, weights = train_classifiers_with_bias(sampled_train_x,
90 sampled_train_y, sampled_test_x, sampled_test_y, λ_val, learning_rate, epochs)
91     push!(new_accuracies_per_learning_rate, accuracies)
92     push!(new_weights_per_learning_rate, weights)
93 end
94
95 # Display overall accuracy for all lambdas
96 println("Overall Accuracies for All learning_rates:")
97 for (learning_rates, accs) in zip(learning_rates,
98 new_accuracies_per_learning_rate)
99     println("learning_rates = ", learning_rates, ": ", mean(accs))
100 end
```


Layer with Float32 parameters got Float64 input.

The input will be converted, but any earlier layers may be very slow.

layer: Dense(785 => 1, σ) # 786 parameters


summary(x): "785x100 Matrix{Float64}"

```
learning_rate = 0.001
Training classifier for digit 0
Accuracy for digit 0: 0.94
Confusion Matrix for digit 0:
[94 0; 6 0]
Training classifier for digit 1
Accuracy for digit 1: 0.86
Confusion Matrix for digit 1:
[85 3; 11 1]
Training classifier for digit 2
Accuracy for digit 2: 0.87
Confusion Matrix for digit 2:
[86 1; 12 1]
Training classifier for digit 3
Accuracy for digit 3: 0.87
Confusion Matrix for digit 3:
[87 0; 13 0]
Training classifier for digit 4
Accuracy for digit 4: 0.87
Confusion Matrix for digit 4:
[87 2; 11 0]
Training classifier for digit 5
Accuracy for digit 5: 0.86
Confusion Matrix for digit 5:
[86 2; 12 0]
Training classifier for digit 6
```



(10 pts) section b: Model analysis

For each of your 10 trained classifiers, show an image of the 784 weights (don't show the bias weight) as a 28×28 image. Do these images provide any insight to how this classifier works?



```

1 begin
2   using Plots: heatmap
3
4
5   # Preprocess data (normalize each image individually)
6   function nw_preprocess_data(images, labels, digit)
7     X = Float32.(images) # Convert to Float32
8     X = reshape(X, :, size(X, 3))
9     X = [(x .- mean(x)) / std(x) for x in eachcol(X)] # Normalize each image
10    individually
11    X = hcat(X...)
12    X = vcat(X, ones(1, size(X, 2)))
13    y = (labels .== digit)
14    return convert(Matrix{Float32}, X), y # Convert X to Matrix{Float32}
15  end
16
17  # Train classifiers for each digit with bias
18  function weight_train_classifiers_with_bias(train_x, train_y, test_x, test_y, λ,
19  lr, epochs)
20    accuracies = []
21    confusion_matrices = []
22    weights_matrices = [] # Store weights matrices for visualization
23    for digit in 0:9 # Binary classification
24      X_train, y_train = nw_preprocess_data(train_x, train_y, digit)
25      X_test, y_test = nw_preprocess_data(test_x, test_y, digit)
26      model = new_logistic_regression_model(size(X_train, 1))
27      new_train_model!(model, X_train, y_train, λ, lr, epochs)
28
29      # Check if there are any weight parameters
30      weights = params(model)
31      if isempty(weights)
32        println("No weight parameters found for classifier $digit")
33        push!(accuracies, NaN)
34        continue
35      end
36    end
37  end

```

```
34
35     ŷ = model(X_test)
36     predictions = Flux.round.(Int, vec(ŷ))
37     accuracy = mean(predictions .== y_test)
38     push!(accuracies, accuracy)
39
40     # Calculate and store confusion matrix
41     confusion_matrix = calculate_confusion_matrix(predictions, y_test)
42     push!(confusion_matrices, confusion_matrix)
43
44     # Store weights for visualization, excluding bias weights
45     push!(weights_matrices, weights[1]) # Exclude bias weights
46 end
47 return accuracies, confusion_matrices, weights_matrices
48 end
49 function plot_weights_per_digit(weights_matrices)
50     plots = []
51     for i in 1:10
52         reshaped_data = reshape(weights_matrices[i][1, 1:784], 28, 28)
53         heatmap_plot = heatmap(reshaped_data, c=:binary, yflip=true, colorbar=false)
54         push!(plots, heatmap_plot)
55     end
56     plot(plots..., layout=(2, 5), size=(1200, 600))
57 end
58 # Train classifiers and store weights for visualization
59 nw_accuracies_per_learning_rate = []
60 nw_weights_per_learning_rate = []
61 for learning_rate in learning_rates
62     accuracies, _, weights = weight_train_classifiers_with_bias(sampled_train_x,
63 sampled_train_y, sampled_test_x, sampled_test_y, λ_val, learning_rate, epochs)
64     push!(nw_accuracies_per_learning_rate, accuracies)
65     push!(nw_weights_per_learning_rate, weights)
66 end
67 plot_weights_per_digit(nw_weights_per_learning_rate[end])
68 end
69
```